

**2017.1 Human Media Multicore Computing**  
**Final Exam (June 16th 10am-11:20am)**

supervisor	
signature	

StudentID# : ( ) , Name : ( )

\* You may answer in either Korean or English language unless instructed to answer in English.

1. (18 points) Fill out the blanks (a)~(i) with the most appropriate English words.
- GPGPU stands for (a. G ) (b. P ) computing on Graphics Processing Units. GPGPU is the use of GPU, which typically handles computation only for graphics, to perform computation in applications traditionally handled by CPU.
  - CUDA C/C++ keyword \_\_global\_\_ indicates a function
    - is executed on (c. ), and
    - is called from (d. ).
 Any call to a \_\_global\_\_ function must specify (e. ) for that call.
  - In GPU, a stream multiprocessor (SM) is basically (f. ) processor that executes a warp simultaneously.
  - [In OpenMP] By default, all variables declared outside a parallel block are (g. ), except (h. ) variable, which is (i. ).

2. (10 points) [In pthread library] Consider a program statement "**pthread\_join(A,B)**". Answer to following questions.

- (1) What does the function **pthread\_join** do? Explain **pthread\_join** with sufficient details. ( )
- (2) What is the purpose of the argument variable **A**? Explain **A** with sufficient details. ( )
- (3) What is the purpose of the argument variable **B**? Explain **B** with sufficient details. ( )

3. (12 points) Following program in the left box intends to compute the sum between 1 and 10000 with multiple threads using OpenMP. However, the program is not correct and may generate a wrong result.

- (1) Why is the program (in the left box) wrong? Explain with sufficient details. ( )
- (2) Insert a correct code in the right empty box.
- (3) Explain how/why your code can make the program correct. ( )

<pre>#include &lt;omp.h&gt; #include &lt;stdio.h&gt; #define NUM_THREADS 4  int main () {     int i, sum=0;     omp_set_num_threads(NUM_THREADS);     #pragma omp parallel for         for (i=1; i&lt;=10000; i++) {             sum += i;         }     printf("sum = 1+2+...+10000 = %d\n", sum);     return 0; }</pre>	<pre>#include &lt;omp.h&gt; #include &lt;stdio.h&gt; #define NUM_THREADS 4  int main () {     // insert your correct code here      printf("sum = 1+2+...+10000 = %d\n", sum);     return 0; }</pre>
---	--

4. (15 points) Following pseudocode, which we learned in our class, describes parallel mergesort algorithm **Par-Merge-Sort** using a divide-and-conquer approach. Assume that the function **Par-Merge** correctly defines the parallel merge algorithm. **spawn** means creating and starting a new thread. Also, assume that we use a typical multicore computer.

(1) What is the most serious problem of the following parallel mergesort algorithm **Par-Merge-Sort**, when we run the algorithm in real software? Explain with sufficient details. ( )

(2) How can you modify the algorithm to solve the problem? Please modify or insert your code directly in the following pseudocode.

```
Par-Merge-Sort ( A, p, r ) { sort the elements in A[ p ... r ] }

1. if p < r then
2.   q ← ⌊ ( p + r ) / 2 ⌋
3.   spawn Par-Merge-Sort ( A, p, q )
4.   Par-Merge-Sort ( A, q + 1, r )
5.   sync
6.   Par-Merge ( A, p, q, r )
```

5. (20 points) Answer to following questions that are related to prefix sum by filling out empty boxes with appropriate pseudocodes.
- (a) In prefix sum algorithm, input is a sequence of  $n$  elements  $\{x_1, x_2, \dots, x_n\}$  with a binary associative operation (binary addition) denoted by  $\oplus$ , and output is  $\{s_1, s_2, \dots, s_n\}$ , where  $s_i =$   for  $1 \leq i \leq n$ .
- (b) Fill out the empty box in the following pseudocode for parallel prefix sum algorithm, which uses divide-and-conquer approach.

-----

**ParallelPrefixSum** ( $\langle x_1, \dots, x_n \rangle, \oplus$ )

if  $n=1$  then

$s_1 \leftarrow x_1;$

else {

}

return  $\langle s_1, \dots, s_n \rangle;$

- 6.(25 points) Consider following CUDA code that multiplies two matrices, M and N, with arbitrary size. (1) Insert appropriate code into empty boxes (a) ~ (h). In the code, BLOCK\_SIZE and WIDTH represents the size of a block and the width of matrices.

```
#include <stdio.h>
#include <sys/time.h>
#define BLOCK_SIZE 32
#define WIDTH 1027

typedef struct {
    int width;
    int height;
    float* elements;
} Matrix;

__global__ void MatrixMulKernel(Matrix M, Matrix N, Matrix P)
{
    (a)
}

Matrix FuncA(const Matrix M) {
    (b)
}

void FuncB(Matrix Mdevice, const Matrix Mhost) {
    (c)
}

void FuncC(Matrix M) {
    (d)
}
```

```
void FuncD(Matrix Mhost, const Matrix Mdevice) {
    (e)
}

void FreeMatrix(Matrix M) { free(M.elements); }

void MatrixMulOnDevice(const Matrix M, const Matrix N, Matrix P) {
    Matrix Md = FuncA(M);
    FuncB(Md, M);
    Matrix Nd = FuncA(N);
    FuncB(Nd, N);
    Matrix Pd = FuncA(P);
    FuncB(Pd, P);

    dim3 dimGrid( (f) , (g) );
    dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);

    MatrixMulKernel (h) ;

    FuncD(P, Pd);

    FuncC(Md); FuncC(Nd); FuncC(Pd);
}

Matrix AllocateMatrix(int height, int width) {
    Matrix M; M.width = width; M.height = height;
    int size = M.width * M.height;
    M.elements = (float*) malloc(size*sizeof(float));
    for (unsigned int i = 0; i < M.height * M.width; i++)
        M.elements[i] = 1.0;
    return M;
}

int main(void) {
    Matrix M = AllocateMatrix(WIDTH, WIDTH);
    Matrix N = AllocateMatrix(WIDTH, WIDTH);
    Matrix P = AllocateMatrix(WIDTH, WIDTH);
    MatrixMulOnDevice(M, N, P);
    cudaDeviceSynchronize();
    FreeMatrix(M); FreeMatrix(N); FreeMatrix(P);
    return 0;
}
```