

Multicore Processor

A single computing component with two or more independent cores

Multiple cores run multiple instructions at the same time (concurrently)

Increase overall program speed (performance)

Manycore processor (GPU)

multi-core architectures with an especially high number of cores

Process

run in separate memory address space

a program in execution

Thread

Run in shared memory space in a process

One process may have multiple threads

Multithreaded Program

A program running with multiple threads that is executed simultaneously

Parallel Computing

Using multiple processors in parallel to solve problems more quickly than with a single processor

All processors may have access to a shared memory to exchange information between processors

More tightly coupled to multi-threading

Parallel Programming

Using additional computational resources to produce an answer faster

Concurrent Programming

Correctly and efficiently controlling access by multiple threads to shared resources

Problem of preventing a bad interleaving of operations from different threads

Distributed Computing

Multiple computers communicate through network

Each processor has its own private memory (distributed memory)

Executing sub-tasks on different machines and then merging the results

Cluster Computing

A set of loosely connected computers that work together so that in many respects they can be viewed as single system

Good price / performance, memory not shared

Grid Computing

Federation of computer resources from multiple locations to reach a common goal
Grids tend to be more loosely coupled, heterogeneous, and geographically dispersed

Moore's Law

Doubling of the number of transistors on integrated circuits roughly every two year

Computer Hardware Trend

No more hidden parallelism to be found (ILP)
Transistor number still rising
Clock speed flattening sharply (power consumption, heat generation)
Minimized wire lengths and interconnect latencies

Symmetric MultiProcessor System

Two or more identical processors are connected to a single shared memory
Most common multiprocessor systems today use an SMP architecture

GPGPU

General Purpose Graphical Processing Unit

Overhead of Parallelism

Cost of starting a thread or process
Cost of communicating shared data
Cost of synchronizing
Extra (redundant) computation

Load Imbalance

Some processors are idle due to insufficient parallelism / unequal size tasks

Single/Multiple Instruction Single/Multiple Data

Decomposition – Assignment – Orchestration – Mapping

Decomposition — coverage

Break up computation into tasks to be divided among processes
Identify concurrency and decide level at which to exploit it
Domain – data / Functional – computation, problem

Assignment — granularity

Assign tasks to threads (static/dynamic)
Balance workload, reduce communication and management cost
Together with decomposition, also called partitioning

Orchestration — locality

Structuring (reduce) communication and synchronization

Organizing data structures in memory (reserve locality) and scheduling tasks temporally

Mapping — locality

Mapping threads to execution units (cores) (OS job)

Place related threads on the same processor

Maximize locality, data sharing, minimize costs of comm/sync

Coverage: Amdahl's Law

Potential program speedup is defined by the fraction of code that can be parallelized

Scalability

Capability of a system to increase total throughput under an increased load when resources (typically hardware) are added

Granularity

Qualitative measure of the ratio of computation to communication

Extent to which a system is broken down into small parts

Coarse grained

Relatively large amounts of computational work done between communication events

Consists of fewer, larger components than fine-grained systems

Regards large subcomponents

High computation to communication ratio

More opportunity for performance increase

Advantageous (overhead = comm+sync)

Fine grained

Relatively small amounts of computational work done between communication events

Regards smaller subcomponents of which the larger ones are composed

Low computation to communication ratio

Less opportunity for performance increase

High communication overhead

Load Balancing

Distributing approximately equal amounts of work among tasks so that all tasks are kept busy all of the time

Minimization of task idle time

Static Load Balancing

Programmer make decision and assign fixed amount of work to each core

Low run time overhead

Homogeneous multicores

Dynamic Load Balancing

When one core finishes its allocated work, it takes work from a work queue or core with the heaviest workload

Adapt partitioning at run time to balance load

High runtime overhead

Ideal for codes where work is uneven, unpredictable, heterogeneous multicore

Communication

With message passing, programmer has to understand the computation and orchestrate the communication accordingly

Point to Point / Broadcast + Reduce / All to All / Scatter + Gather

Inter-task communication virtually always implies overhead

Require some type of synchronization between tasks -> waiting

Latency vs Bandwidth / Synchronous vs Asynchronous / Blocking vs Non-blocking

Message Passing Library

Not a language or compiler specification, specific implementation or product

Synchronization

Coordination of simultaneous events (threads / processes) in order to obtain correct runtime order and avoid unexpected condition

Barrier

Any thread/process must stop at this point and cannot proceed until all other threads/processes reach this barrier

Lock / Semaphore

First task acquires the lock – safely access the protected data/code

Other tasks attempt to acquire lock but must wait until the task that owns the lock releases it

Non-Uniform Access

Physically partitioned but accessible by all

Same address space

Placement of data affects performance

Cache-Coherent NUMA

Cache Coherence

Uniformity of shared resource data that ends up stored in multiple local caches

When processor modifies shared variable in local cache, different processors may have different value -> need to take actions to ensure visibility or cache coherence

Shared Memory Architecture

- Global address space provides user-friendly programming perspective to memory
- Fast and uniform data sharing between tasks (due to proximity of memory to CPUs)
- Lack of scalability between memory and CPUs
- Programmer responsibility for synchronization
- Difficult and expensive to design/produce

Distributed Memory Architecture

- Scalable (processors, memory)
- Cost effective
- Programmer responsibility of data communication
- No global memory access
- Non-uniform memory access time

Thread

- Single sequential flow of control within a program
- Program counter + Register set + Stack
- Share - Code section + Data section + OS resources (file)

yield()

release the right of cpu, allows the scheduler to select another runnable thread

join()

one thread can wait for another thread to end

wait()

current thread blocked, placed into wait set, lock object released

notify()

one thread removed from wait set, retain lock for object, resumed from waiting status

Condition Variable

- Synchronization primitives that enable threads to wait until a particular condition occurs
- Enable threads to atomically release a lock and enter the sleeping state
- Without condition variables, programmer need to have threads continually polling
- Always used in conjunction with mutex lock

Producer-Consumer Problem

- Producer thread do some work and enqueue result objects
- Consumer thread dequeue objects and de next stage
- Must synchronize access to queue

Mutual Exclusion

- Prevent more than one thread from accessing critical section at given time